

Class07: Machine Learning 1

Yang Liu

Today we will start out multi-part exploration of some key machine learning methods. We will begin with clustering--finding groupings in data, and then dimensionality reduction.


##Clustering

Let's start with "k-means" clustering


The main function in base R for this 'kmeans()'.

```
#make up some data  
hist(rnorm(100000, mean=3))
```

Histogram of rnorm(1e+05, mean = 3)



```
tmp <- c(rnorm(30, -3), rnorm(30, +3))
x <- cbind(x=tmp, y=rev(tmp))
plot(x)
```



Now let's try out 'kmeans()'

```
km <- kmeans(x, centers=2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	3.142578	-3.190014
2	-3.190014	3.142578

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

Available components:

```
attributes(km)
```

```
$class
[1] "kmeans"
```

km\$size


Q. What component of your result object details cluster assignment/membership?

[illegible]

km\$centers

Q. Make a plot of your data showing your clustering results (grouping/clusters and cluster centers).

```
plot(x, col=c("red","blue"))
```




```
c(1:5) + c(100, 1)
```


Warning in `c(1:5) + c(100, 1)`: longer object length is not a multiple of
shorter object length

```
[1] 101    3 103    5 105
```

```
plot(x, col=c(1,2))
```



```
plot(x, col=km$cluster)
points(km$centers, col="green", pch=15, cex=3)
```



Q. Run 'kmeans' again and cluster in 4 groups and plot the results.

```
km4 <- kmeans(x, centers=4)
km4
```

K-means clustering with 4 clusters of sizes 14, 9, 16, 21

Cluster means:

	x	y
1	-3.755358	3.713319
2	4.222433	-3.642528
3	-2.695338	2.643180
4	2.679783	-2.996080

Clustering vector:

[1] 1 1 1 3 3 3 3 3 3 3 3 1 1 1 3 1 1 1 3 1 3 3 1 3 3 1 3 3 1 1 2 2 4 4 4 4 4 2
[39] 4 4 2 4 4 2 2 4 2 2 2 4 4 4 4 4 4 4 4 4 4 4


Within cluster sum of squares by cluster:

```
[1] 12.634305  3.812884 13.248033 21.384833
(between_SS / total_SS = 96.0 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
plot(x, col=km4$cluster)
points(km4$centers, col="green", pch=15, cex=3)
```



##Hierarchical Clustering

This form of clustering aims to reveal the structure in your data by progressively grouping points into a ever smaller number of clusters.

The main fucntion in base R for this called ‘hclust()’ This function does not take out input data directiely but wants a “distance matrix” that details how (dis)similar all our inout points are to each otehr.

```
hc <- hclust(dist(x))
hc
```

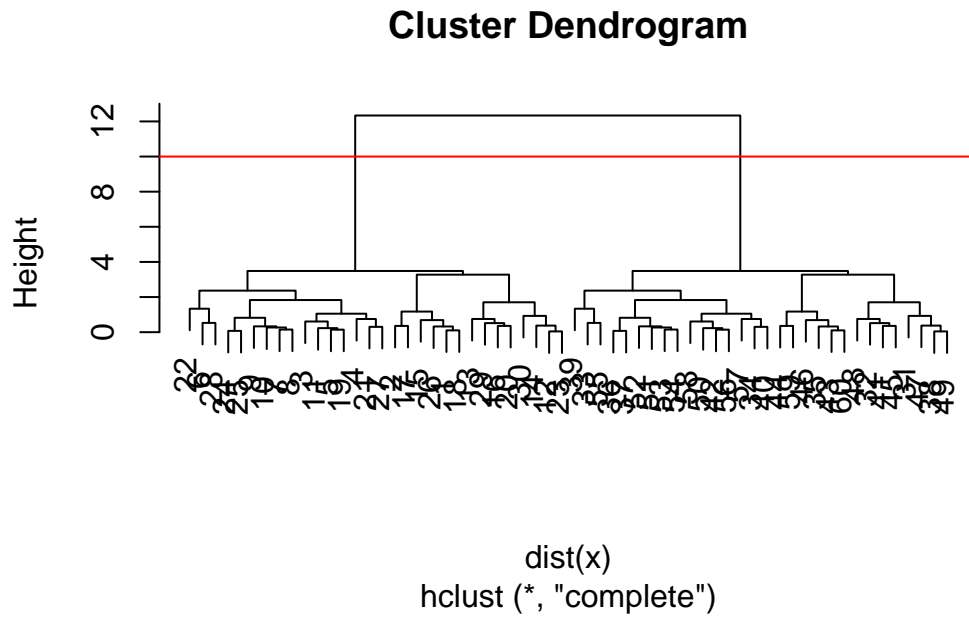
Call:

```
hclust(d = dist(x))
```

```
Cluster method   : complete
Distance         : euclidean
Number of objects: 60
```

The print out above is not very useful (unlick that from kmeans) but there is a useful ‘plot()’ mehtod.

```
plot(hc)
abline(h=10, col="red")
```




To get my main result (My cluster membership vector) I need to “cut” my tree using the function ‘cutree()’


```
grps <- cutree(hc, h=10)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
plot(x, col=grps)
```

```
plot(x, col=cutree(hc, h=6))
```



#Principal Component Analysis (PCA)

The goal of PCA is to reduce the dimension of a dataset down to some smaller subset of new variables (called PCs) that are a useful based for furthur analysis, like visualization, clustering, etc.

Data import

Read data about crazy eacting trends in the UK and N. Ireland.

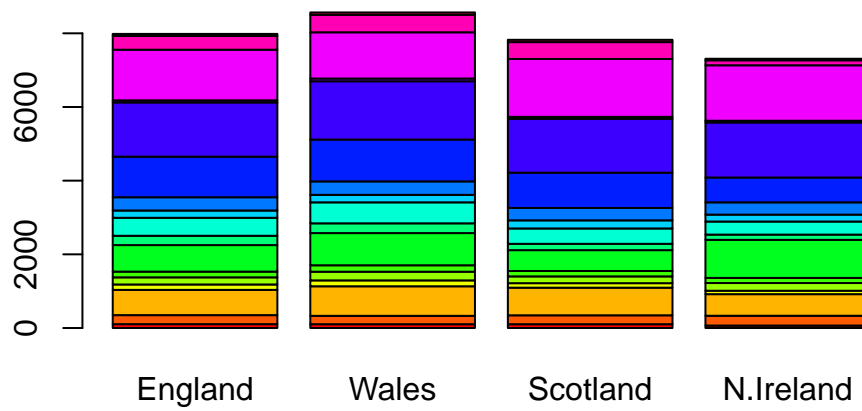
```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41


```
dim(x)
```

```
[1] 17  4
```

```
barplot(as.matrix(x), col=rainbow(nrow(x)))
```




```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



The so-called “pairs” plot can be useful for small datasets:

```
#rainbow(nrow(x))
pairs(x, col=rainbow(nrow(x)), pch=20)
```



So the paris plot is useful for small datasets but it can be lots of work to interpret and get untractable for larger datasets.

So PCA to the rescue...

The main function to do PCA in base R is called ‘prcomp()’ This function wants the transpose of our data in this case

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```


```
[1] "prcomp"
```

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

A major PCA result viz is called a “PCA plot” (a.k.a. a score plot, biplot, PC1 vs. PC2 plot, ordination plot)

```
mycols <- c("orange", "red", "blue", "darkgreen")
plot(pca$x[, 1], pca$x[,2], col=mycols, pch=16,
     xlab="PC1", ylab="PC2" )
abline(h=0, col="gray")
abline(v=0, col="gray")
```



Another important output from PCA is called the “loadings” vector or the “rotation” component -this tell us how much the original variables (the food in this case) contributes to the new PCs

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.694538519
Carcass_meat	0.047927628	0.013915823	0.06367111	0.489884628
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.279023718
Fish	-0.084414983	-0.050754947	0.03906481	-0.008483145
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.076097502
Sugars	-0.037620983	-0.043021699	-0.03605745	0.034101334
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	-0.090972715
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	-0.039901917
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.016719075
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	0.030125166
Processed_Veg	-0.036488269	-0.045451802	0.05289191	-0.013969507
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.184072217
Cereals	-0.047702858	-0.212599678	-0.35884921	0.191926714
Beverages	-0.026187756	-0.030560542	-0.04135860	0.004831876
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.103508492

Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.316290619
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001847469

PCA looks to be a super useful method for gaining some insight into high dimensional data that is difficult to examine in other ways.

PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638